

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import datetime as dt
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

# Load dataset
df = pd.read_csv('sales_data_sample.csv', encoding='unicode_escape')

df.head()

```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	
SALES \					
0	10107	30	95.70	2	2871.00
1	10121	34	81.35	5	2765.90
2	10134	41	94.74	2	3884.34
3	10145	45	83.26	6	3746.70
4	10159	49	100.00	14	5205.27

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	... \
0	2/24/2003 0:00	Shipped	1	2	2003	...
1	5/7/2003 0:00	Shipped	2	5	2003	...
2	7/1/2003 0:00	Shipped	3	7	2003	...
3	8/25/2003 0:00	Shipped	3	8	2003	...
4	10/10/2003 0:00	Shipped	4	10	2003	...

	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME	
DEALSIZE						
0	10022	USA	NaN	Yu	Kwai	
Small						
1	51100	France	EMEA	Henriot	Paul	
Small						
2	75508	France	EMEA	Da Cunha	Daniel	
Medium						
3	90003	USA	NaN	Young	Julie	
Medium						
4	NaN	USA	NaN	Brown	Julie	
Medium						

```
[5 rows x 25 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ORDERNUMBER      2823 non-null    int64  
 1   QUANTITYORDERED 2823 non-null    int64  
 2   PRICEEACH        2823 non-null    float64 
 3   ORDERLINENUMBER 2823 non-null    int64  
 4   SALES            2823 non-null    float64 
 5   ORDERDATE        2823 non-null    object  
 6   STATUS            2823 non-null    object  
 7   QTR_ID           2823 non-null    int64  
 8   MONTH_ID         2823 non-null    int64  
 9   YEAR_ID          2823 non-null    int64  
 10  PRODUCTLINE      2823 non-null    object  
 11  MSRP              2823 non-null    int64  
 12  PRODUCTCODE      2823 non-null    object  
 13  CUSTOMERNAME     2823 non-null    object  
 14  PHONE             2823 non-null    object  
 15  ADDRESSLINE1     2823 non-null    object  
 16  ADDRESSLINE2     302 non-null     object  
 17  CITY              2823 non-null    object  
 18  STATE             1337 non-null    object  
 19  POSTALCODE        2747 non-null    object  
 20  COUNTRY           2823 non-null    object  
 21  TERRITORY         1749 non-null    object  
 22  CONTACTLASTNAME   2823 non-null    object  
 23  CONTACTFIRSTNAME  2823 non-null    object  
 24  DEALSIZE          2823 non-null    object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB

# Drop unnecessary columns
to_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATE', 'POSTALCODE',
           'PHONE']
df = df.drop(to_drop, axis=1)

#Check for null values
df.isnull().sum()

ORDERNUMBER      0
QUANTITYORDERED 0
PRICEEACH        0
ORDERLINENUMBER  0
```

```

SALES          0
ORDERDATE      0
STATUS          0
QTR_ID          0
MONTH_ID        0
YEAR_ID         0
PRODUCTLINE     0
MSRP            0
PRODUCTCODE     0
CUSTOMERNAME    0
CITY             0
COUNTRY          0
TERRITORY       1074
CONTACTLASTNAME 0
CONTACTFIRSTNAME 0
DEALSIZE         0
dtype: int64

```

*#Bhai bhai look at territory, Vishwajeet shocked other admin rocked
#But territory does not have significant impact on analysis like RFM segmentation or clustering so ignore it*

df.dtypes

ORDERNUMBER	int64
QUANTITYORDERED	int64
PRICEEACH	float64
ORDERLINENUMBER	int64
SALES	float64
ORDERDATE	object
STATUS	object
QTR_ID	int64
MONTH_ID	int64
YEAR_ID	int64
PRODUCTLINE	object
MSRP	int64
PRODUCTCODE	object
CUSTOMERNAME	object
CITY	object
COUNTRY	object
TERRITORY	object
CONTACTLASTNAME	object
CONTACTFIRSTNAME	object
DEALSIZE	object
dtype:	object

#ORDERDATE Should be in date time

```
df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])
```

```

# RFM Feature Calculation
snapshot_date = df['ORDERDATE'].max() + dt.timedelta(days=1)
df_RFМ = df.groupby('CUSTOMERNAME').agg({
    'ORDERDATE': lambda x: (snapshot_date - x.max()).days,
    'ORDERNUMBER': 'count',
    'SALES': 'sum'
})

# Rename columns
df_RFМ.rename(columns={
    'ORDERDATE': 'Recency',
    'ORDERNUMBER': 'Frequency',
    'SALES': 'MonetaryValue'
}, inplace=True)

# Quartile segmentation
df_RFМ['M'] = pd.qcut(df_RFМ['MonetaryValue'], q=4, labels=range(1,5))
df_RFМ['R'] = pd.qcut(df_RFМ['Recency'], q=4, labels=list(range(4,0,-1)))
df_RFМ['F'] = pd.qcut(df_RFМ['Frequency'], q=4, labels=range(1,5))

df_RFМ.head()

```

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F
AV Stores, Co.	196	51	157807.81	4	2	4
Alpha Cognac	65	20	70488.44	2	4	2
Amica Models & Co.	265	26	94117.26	3	1	2
Anna's Decorations, Ltd	84	46	153996.13	4	3	4
Atelier graphique	188	7	24179.96	1	2	1

```

# RFM Score
df_RFМ['RFM_Score'] = df_RFМ[['R', 'F', 'M']].sum(axis=1)

```

```

# RFM Level
def rfm_level(df):
    if df['RFM_Score'] >= 10:
        return 'High Value Customer'
    elif 6 <= df['RFM_Score'] < 10:
        return 'Mid Value Customer'
    else:
        return 'Low Value Customer'

```

```

df_RFМ['RFM_Level'] = df_RFМ.apply(rfm_level, axis=1)
print(df_RFМ.head())

```

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F	\
AV Stores, Co.	196	51	157807.81	4	2	4	
Alpha Cognac	65	20	70488.44	2	4	2	
Amica Models & Co.	265	26	94117.26	3	1	2	

Anna's Decorations, Ltd	84	46	153996.13	4	3	4
Atelier graphique	188	7	24179.96	1	2	1

CUSTOMERNAME	RFM_Score	RFM_Level
AV Stores, Co.	10	High Value Customer
Alpha Cognac	8	Mid Value Customer
Amica Models & Co.	6	Mid Value Customer
Anna's Decorations, Ltd	11	High Value Customer
Atelier graphique	4	Low Value Customer

```
# Data Preparation for Clustering
data = df_RFU[['Recency', 'Frequency', 'MonetaryValue']]
```

We create levels for our Customers

RFM Score > 10 : High Value Customers

RFM Score < 10 and RFM Score >= 6 : Mid Value Customers

RFM Score < 6 : Low Value Customers

```
# Log transformation to reduce skew
data_log = np.log1p(data) # log(1 + x)

# Standardization
scaler = StandardScaler()
data_normalized = pd.DataFrame(scaler.fit_transform(data_log),
                                index=data_log.index,
                                columns=data_log.columns)
print(data_normalized.describe().round(2))
```

	Recency	Frequency	MonetaryValue
count	92.00	92.00	92.00
mean	0.00	0.00	0.00
std	1.01	1.01	1.01
min	-3.26	-3.41	-3.82
25%	-0.29	-0.43	-0.39
50%	0.37	0.05	-0.04
75%	0.54	0.44	0.52
max	1.18	4.16	3.92

```
# Our data is skewed we must remove it by performing log
transformation
data_log = np.log(data)
data_log.head()
```

CUSTOMERNAME	Recency	Frequency	MonetaryValue
AV Stores, Co.	5.278115	3.931826	11.969133
Alpha Cognac	4.174387	2.995732	11.163204
Amica Models & Co.	5.579730	3.258097	11.452297
Anna's Decorations, Ltd	4.430817	3.828641	11.944683
Atelier graphique	5.236442	1.945910	10.093279

```
#Fit KMeans and use elbow method to choose the number of clusters
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans

sse = {}

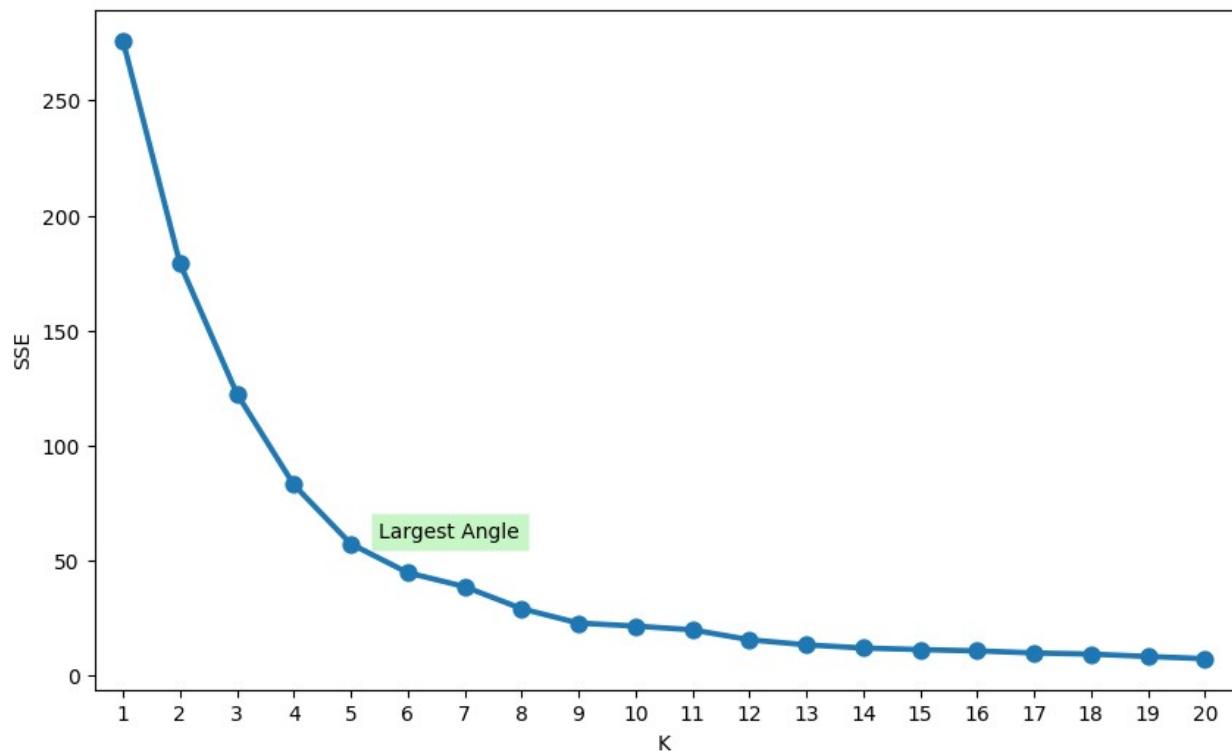
for k in range(1, 21):
    kmeans = KMeans(n_clusters = k, random_state = 1)
    kmeans.fit(data_normalized)
    sse[k] = kmeans.inertia_

plt.figure(figsize=(10,6))
plt.title('The Elbow Method')

plt.xlabel('K')
plt.ylabel('SSE')
plt.style.use('ggplot')

sns.pointplot(x=list(sse.keys()), y = list(sse.values()))
plt.text(4.5, 60, "Largest Angle", bbox = dict(facecolor = 'lightgreen', alpha = 0.5))
plt.show()
```

The Elbow Method



```
kmeans = KMeans(n_clusters=5, random_state=1)
kmeans.fit(data_normalized)
cluster_labels = kmeans.labels_

data_rfm = data.assign(Cluster = cluster_labels)
data_rfm.head()
```

CUSTOMERNAME	Recency	Frequency	MonetaryValue	Cluster
AV Stores, Co.	196	51	157807.81	4
Alpha Cognac	65	20	70488.44	2
Amica Models & Co.	265	26	94117.26	2
Anna's Decorations, Ltd	84	46	153996.13	4
Atelier graphique	188	7	24179.96	1