

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier

# Load dataset
df = pd.read_csv("emails.csv")

df.head()

   Email No.    the    to    ect    and    for    of     a    you    hou    ...    connevey
jay \
0   Email 1    0    0     1    0     0    0     2    0    0    0    ...
0
1   Email 2    8   13    24    6     6    2   102    1   27    ...
0
2   Email 3    0    0     1    0     0    0     8    0    0    ...
0
3   Email 4    0    5    22    0     5    1   51    2   10    ...
0
4   Email 5    7    6    17    1     5    2   57    0    9    ...
0

   valued    lay  infrastructure  military  allowing    ff    dry
Prediction
0       0    0                 0        0        0    0    0
0
1       0    0                 0        0        0    1    0
0
2       0    0                 0        0        0    0    0
0
3       0    0                 0        0        0    0    0
0
4       0    0                 0        0        0    1    0
0

[5 rows x 3002 columns]

df.isnull().sum()

Email No.      0
the            0
to             0
ect            0
and            0
...
military      0
allowing      0
ff             0
dry            0

```

```

Prediction      0
Length: 3002, dtype: int64

X = df.iloc[:, 1:3001] # word frequency features
X

      the   to   ect   and   for   of     a   you   hou   in   ...   enhancements
\0      0     0     1     0     0     0     2     0     0     0     0     ...
1       8    13    24    6     6     2    102    1    27    18    ...
2       0     0     1     0     0     0     8     0     0     4     ...
3       0     5    22    0     5     1    51    2    10    1     ...
4       7     6    17    1     5     2    57    0     9     3     ...
...
5167    2     2     2     3     0     0    32    0     0     5     ...
5168    35    27    11    2     6     5   151    4     3    23    ...
5169    0     0     1     1     0     0    11    0     0     1     ...
5170    2     7     1     0     2     1    28    2     0     8     ...
5171   22    24     5     1     6     5   148    8     2    23    ...

      connevey   jay   valued   lay   infrastructure   military   allowing
ff   dry
0       0     0     0     0           0           0           0
0       0
1       0     0     0     0           0           0           0
1       0
2       0     0     0     0           0           0           0
0       0
3       0     0     0     0           0           0           0
0       0
4       0     0     0     0           0           0           0
1       0
...
5167    0     0     0     0           0           0           0
0       0
5168    0     0     0     0           0           0           0
1       0
5169    0     0     0     0           0           0           0
0       0

```

```

5170      0  0  0  0          0  0  0
1      0
5171      0  0  0  0          0  0  0
0      0

[5172 rows x 3000 columns]

Y = df.iloc[:, -1].values # 1 = spam, 0 = not spam
Y

array([0, 0, 0, ..., 1, 1, 0], shape=(5172,))

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.25, random_state=42)

from sklearn.metrics import classification_report, confusion_matrix

# ----- Support Vector Machine -----
svc = SVC(C=1.0, kernel='rbf', gamma='auto')
svc.fit(X_train, y_train)
svc_pred = svc.predict(X_test)

SVM Accuracy: 0.8932714617169374
SVM Classification Report:
precision    recall   f1-score   support
      0       0.90      0.96      0.93      913
      1       0.87      0.74      0.80      380
      accuracy           0.89
      macro avg       0.89      0.85      0.87      1293
      weighted avg     0.89      0.89      0.89      1293

SVM Confusion Matrix:
[[872  41]
 [ 97 283]]

print("SVM Accuracy:", accuracy_score(y_test, svc_pred))
print("SVM Classification Report:\n", classification_report(y_test,
svc_pred))
print("SVM Confusion Matrix:\n", confusion_matrix(y_test, svc_pred))

SVM Accuracy: 0.8932714617169374
SVM Classification Report:
precision    recall   f1-score   support
      0       0.90      0.96      0.93      913
      1       0.87      0.74      0.80      380
      accuracy           0.89
      1293

```

| | | | | |
|--------------|------|------|------|------|
| macro avg | 0.89 | 0.85 | 0.87 | 1293 |
| weighted avg | 0.89 | 0.89 | 0.89 | 1293 |

SVM Confusion Matrix:

```
[[872 41]
 [ 97 283]]
```

----- K-Nearest Neighbors -----

```
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)

print("KNN Accuracy:", knn.score(X_test, y_test))
print("KNN Classification Report:\n", classification_report(y_test,
knn_pred))
print("KNN Confusion Matrix:\n", confusion_matrix(y_test, knn_pred))
```

KNN Accuracy: 0.8685990338164251

KNN Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.87 | 0.90 | 739 |
| 1 | 0.73 | 0.86 | 0.79 | 296 |
| accuracy | | | 0.87 | 1035 |
| macro avg | 0.83 | 0.87 | 0.85 | 1035 |
| weighted avg | 0.88 | 0.87 | 0.87 | 1035 |

KNN Confusion Matrix:

```
[[645 94]
 [ 42 254]]
```